
Does Cost-Efficiency Travel? Accuracy and Cost Rank Transfer in Fixed-Scaffold Agent Evaluation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agent leaderboards increasingly report both task success and dollar cost, inviting
2 practitioners to select a “cost-efficient” model from a single benchmark. We test
3 whether this selection signal transfers across workloads. Using public results from
4 the Holistic Agent Leaderboard (HAL), we hold the displayed agent scaffold fixed
5 to the HAL Generalist Agent and compare repeated displayed model labels across
6 six benchmarks. Accuracy-rank association is often positive across benchmark
7 pairs, while dollar-cost-rank association is markedly more heterogeneous: among
8 pairs with at least 12 shared labels, cost Spearman correlation ranges from 0.74 to
9 0.93; across all eligible pairs it ranges from -0.75 to 0.93. For example, GAIA
10 and ScienceAgentBench have accuracy correlation 0.64 but cost correlation -0.75
11 over seven shared labels. Correspondingly, nondominated accuracy–cost frontier
12 sets have low overlap. We also distinguish the ordinary nondominated frontier from
13 a HAL-inspired convex-hull reconstruction, finding that the public HAL labels
14 align closely with the latter but not identically. These descriptive results suggest
15 that dollar-cost efficiency from one agent leaderboard should not be presumed
16 portable to another workload, even when accuracy ordering is partly preserved.
17 The study does not identify causal model or scaffold effects: public display labels
18 do not fully encode all benchmark-specific setup details.

19 1 Introduction

20 Agent benchmarks now evaluate systems that reason over multiple steps, invoke tools, and interact
21 with environments. Unlike static model evaluations, an agent’s reported behavior depends on
22 an underlying language model, an agent scaffold, tools, prompts, execution limits, and the task
23 environment. These systems can also be expensive to evaluate: the Holistic Agent Leaderboard (HAL)
24 reports 21,730 rollouts across nine benchmarks and roughly \$40,000 in evaluation cost [Kapoor et al.,
25 2026]. Consequently, cost-aware leaderboards increasingly plot accuracy against dollar cost and
26 recommend points on a cost–accuracy frontier.

27 A practical inference often follows: if a model or agent is cost-efficient on one benchmark, it is a
28 promising low-cost choice elsewhere. This inference is plausible if model capability and model cost
29 are mostly portable. It can fail if an agent’s token consumption, retry loops, tool calls, and stopping
30 behavior change with the workload. This paper asks a narrow descriptive question: *when the publicly
31 displayed scaffold is held fixed, how stable are repeated HAL displayed model labels’ accuracy ranks,
32 dollar-cost ranks, and cost–accuracy frontier positions across workloads?*

33 We conduct a secondary analysis of public HAL leaderboard data. The data are sparse and unbalanced
34 across scaffolds, so rather than attempt an unsupported causal decomposition of model, scaffold, and
35 workload effects, we focus on the only broadly repeated scaffold: the HAL Generalist Agent. Across

36 six benchmarks with sufficient overlap, we compare rankings only among shared displayed model
37 labels. Our main findings are:

- 38 • Accuracy-rank associations are often positive, but dollar-cost-rank associations are substantially
39 more heterogeneous and can be negative.
- 40 • Nondominated cost–accuracy frontier membership is correspondingly unstable: no displayed label
41 tested on at least five of the six workloads is nondominated on every tested workload.
- 42 • Frontier conclusions depend on the decision model. The public HAL Is Pareto labels closely
43 match an origin-anchored convex-hull reconstruction, whereas standard nondominance retains
44 additional points.

45 The result is not universal non-transfer. Some pairs show high agreement in both accuracy and dollar
46 cost. Rather, transfer cannot safely be assumed from a single leaderboard.

47 2 Related work

48 **Cost-aware agent evaluation.** Kapoor et al. [2026] introduced HAL as a standardized, cost-aware
49 agent leaderboard spanning coding, web, science, and customer-service settings. HAL reports dollar-
50 and token-cost Pareto curves and documents substantial scaffold sensitivity. Kapoor et al. [2024]
51 argue that agent evaluation should jointly consider accuracy and cost, using Pareto plots to expose
52 avoidable expenditure. SWE-Effi studies resource-constrained software-engineering agents and
53 reports that effectiveness depends on scaffold–model integration [Fan et al., 2025]. These works
54 establish the importance of cost and scaffold effects, but do not systematically measure whether a
55 cost–accuracy ordering transfers between distinct workloads under a fixed scaffold.

56 **Stability of agent and benchmark rankings.** Ndzomga [2026] ask whether smaller task subsets
57 can preserve *within-benchmark* agent rankings under scaffold and temporal shift. Their target is
58 evaluation reduction, not cross-workload transfer of cost–accuracy trade-offs. In the static-LLM set-
59 ting, Benchmark² proposes cross-benchmark ranking consistency based on Kendall correlation [Qian
60 et al., 2026]. We adapt the cross-benchmark ranking perspective to interactive agents and separately
61 analyze accuracy ranks, dollar-cost ranks, and frontier-set overlap.

62 3 Data, scope, and coverage

63 We use the public `all_leaderboards_costs_HAL.csv` snapshot distributed with Ndzomga [2026].
64 The frozen file contains 242 displayed evaluations across nine benchmarks, 11 displayed scaffolds,
65 and 29 displayed model labels. Accuracy and total dollar cost are complete, but 235 rows report a
66 single run and only seven report two runs. The data therefore support descriptive rank and frontier
67 analysis, not rollout-level uncertainty estimation.

68 Coverage determines the design. Most scaffolds occur on only one benchmark; the HAL Generalist
69 Agent occurs on seven. We retain the six Generalist benchmarks with adequate overlap: CORE-Bench
70 Hard (21 labels), GAIA (17), SciCode (9), ScienceAgentBench (7), SWE-bench Verified Mini (18),
71 and TAU-bench Airline (14). We exclude USACO because it contains one Generalist row. Figure 1
72 shows the resulting incomplete label-by-benchmark matrix.

73 The repeated unit is an exact public Models display string within the exact public Agent Name
74 display string. This includes the visible date and reasoning suffix when present. It does *not* fully
75 specify all benchmark-specific prompts, tools, budgets, or harness settings. Thus, the analysis
76 describes repeated public HAL labels under a repeated displayed scaffold; it is not a controlled
77 base-model experiment.

78 4 Method

79 **Rank transfer.** For each benchmark pair with at least five shared displayed labels, we compute
80 Spearman rank correlation ρ and Kendall’s τ_b separately for accuracy and total dollar cost. Let L_{ij}
81 be labels observed in both benchmarks i and j . We correlate their displayed accuracy values and,
82 separately, their displayed total costs. We report $|L_{ij}|$ for every pair.




Figure 1: Coverage of the primary cohort. A filled cell indicates that the public HAL table contains the displayed model label under the displayed HAL Generalist Agent scaffold. Sparse coverage motivates pairwise analyses restricted to shared labels.

83 To characterize sensitivity to the finite label set, we use 5,000 percentile *configuration-resampling*
84 draws: resample shared labels with replacement and recompute the correlation. These intervals are
85 neither rollout-level uncertainty intervals nor population-model confidence intervals. Degenerate
86 draws without rank variation are excluded and counted. We do not use interval inclusion of zero for
87 hypothesis testing.

88 **Two frontier definitions.** For lower cost and higher accuracy, the nondominated frontier includes
89 point x if no observed point y has cost no greater than x , accuracy at least that of x , and at least one
90 strict improvement. This corresponds to choosing one displayed configuration. HAL’s public analysis
91 code instead adds the origin and computes an upper cost–accuracy hull, which has a randomized-
92 policy interpretation: a mixture of two systems can dominate an interior point in expectation. We
93 implement a *HAL-inspired convex-hull reconstruction* based on the public reference code, while
94 avoiding a claim of exact identity because two supplied labels differ in this frozen snapshot. For both
95 definitions, we report pairwise Jaccard similarity of frontier sets. We distinguish full-cohort frontiers
96 from frontiers recomputed only on a pair’s shared labels.

97 **Scaffold sensitivity.** As a secondary descriptive check, we compare shared displayed labels across
98 scaffolds *within* a benchmark, reporting paired accuracy differences and log cost ratios. This confirms
99 prior scaffold sensitivity rather than estimating a general scaffold effect.

100 5 Results

101 5.1 Accuracy ranks transfer more consistently than dollar-cost ranks

102 Figure 2 is the main result. Accuracy-rank correlations are often positive, while cost-rank correlations
103 vary from strongly positive to strongly negative. GAIA–SWE-bench Mini has 17 shared labels and




Figure 2: Cross-workload rank transfer under the displayed HAL Generalist Agent scaffold. Each off-diagonal cell reports Spearman ρ and the number of shared displayed labels. Accuracy association is often positive, whereas dollar-cost association is more heterogeneous; negative values occur for several pairs involving ScienceAgentBench.

Table 1: Higher-overlap benchmark pairs. Brackets are 5,000-draw configuration-resampling sensitivity intervals, not rollout-level confidence intervals. ND denotes nondominated frontier membership recomputed on the shared-label cohort.

Pair	N	Accuracy ρ	Cost ρ	ND Jaccard
CORE-Bench Hard–GAIA	14	0.60 [0.16, 0.81]	0.93 [0.72, 0.99]	0.43
CORE-Bench Hard–SWE-bench Mini	14	0.48 [-0.05, 0.82]	0.74 [0.29, 0.93]	0.25
CORE-Bench Hard–TAU Airline	12	0.50 [-0.12, 0.86]	0.78 [0.26, 0.98]	0.17
GAIA–SWE-bench Mini	17	0.71 [0.31, 0.91]	0.89 [0.61, 0.99]	0.33
GAIA–TAU Airline	14	0.58 [0.14, 0.85]	0.86 [0.58, 0.97]	0.29
SWE-bench Mini–TAU Airline	14	0.76 [0.40, 0.91]	0.82 [0.46, 0.97]	0.43

shows positive association in both accuracy ($\rho = 0.71$) and cost ($\rho = 0.89$). In contrast, GAIA–ScienceAgentBench has $\rho = 0.64$ for accuracy but $\rho = -0.75$ for dollar cost over seven shared labels. SciCode–SWE-bench Mini is near zero for both metrics ($\rho = 0.15$ accuracy, 0.02 cost, $N = 9$), but its broad configuration-resampling intervals make it a descriptive example rather than precise evidence of no association.

Table 1 lists the six pairs with at least 12 shared labels. The strongest evidence comes from these larger-overlap comparisons: cost ranks are highly aligned in these pairs, but nondominated frontier Jaccard similarity remains low to moderate (0.17–0.43). Thus, even where two scalar ranks agree, binary frontier composition can change as the achievable accuracy–cost trade-off changes.

5.2 Frontier membership is not portable

Figure 3 displays nondominated membership by workload. Among labels tested on at least five workloads, Claude-3.7 Sonnet High has the highest nondominated appearance count, four of six; no broadly tested label is nondominated everywhere. Table 2 summarizes the broadly covered labels. Because membership is discontinuous near a frontier, we interpret this result together with the continuous rank results rather than as a standalone universal ranking.

5.3 Frontier definition and scaffold comparisons matter

The source HAL labels recover 30 frontier rows under the hull reconstruction, with two discrepancies; standard nondominance includes 19 additional non-dominated rows. This is substantively relevant: a discrete buyer choosing one agent and a deployer willing to randomize between agents solve different decision problems.




Figure 3: Nondominated cost–accuracy membership across primary workloads. The final column is each label’s rate over its available workloads. No label tested on at least five workloads is nondominated on all of them.

Table 2: Nondominated frontier appearances for labels tested on at least five primary workloads.

Model label	Tested	ND appearances	ND rate
Claude-3.7 Sonnet High	6	4	0.67
Claude-3.7 Sonnet	6	3	0.50
o4-mini High	6	3	0.50
o4-mini Low	6	3	0.50
DeepSeek V3	6	1	0.17
o3 Medium	6	1	0.17
GPT-4.1	6	0	0.00
Gemini 2.0 Flash	5	3	0.60
DeepSeek R1	5	0	0.00

Within individual benchmarks, scaffold changes are also large and direction-dependent. For example, across 15 shared labels on SWE-bench Mini, the displayed SWE-Agent is on average 27.4 percentage points more accurate than the displayed Generalist scaffold, with a higher median cost. Across seven shared labels on ScienceAgentBench, SAB Self-Debug is both more accurate on average and lower cost. These benchmark-specific contrasts are consistent with HAL and SWE-Effi, and motivate rather than replace the fixed-scaffold primary design.

6 Discussion

The practical implication is conditional, not absolute. A cost-efficient point on one benchmark can be a useful prior, especially when benchmark pairs show strong cost-rank agreement. Yet cost rank is not uniformly portable, and several pairs exhibit weak or negative association despite positive accuracy association. Dollar cost plausibly responds to workload-specific execution dynamics—tool calls, retries, loop length, and termination behavior—in addition to model pricing. A practitioner

136 selecting an agent for a new workload should therefore validate cost on representative tasks rather
137 than extrapolate from a single efficiency leaderboard.

138 The analysis also highlights a reporting issue. A frontier label is not self-explanatory: ordinary non-
139 dominance and a convex-hull/randomization frontier encode different deployment choices. Leader-
140 boards should state which decision model they use and, when feasible, expose both.

141 7 Limitations

142 This is a secondary analysis of a sparse public snapshot. First, repeated coverage is incomplete;
143 most scaffolds occur on only one benchmark, and the primary fixed-scaffold analysis retains only
144 six workloads. Second, public display labels do not fully encode benchmark-specific prompts, tool
145 policies, budgets, or harness versions. Holding a display label fixed does not isolate a causal base-
146 model effect. Third, most evaluations are single-run, so configuration-resampling intervals do not
147 measure rollout variance. Fourth, dollar cost combines observed execution with pricing assumptions;
148 the supplied CSV lacks token totals, preventing a token-cost robustness analysis. Fifth, only three
149 same-domain pairs are available under HAL’s categories, so domain conclusions are exploratory and
150 omitted from the main claim. Finally, frontier membership is sensitive to the comparator set and
151 small changes near the frontier; we therefore report both full-cohort and shared-cohort analyses in
152 the supplemental material.

153 8 Conclusion

154 Across public HAL results with a repeated displayed Generalist scaffold, accuracy ranks often transfer
155 more consistently than dollar-cost ranks, while cost–accuracy frontier membership is not reliably
156 portable. The evidence is heterogeneous rather than uniformly negative: some workload pairs align
157 strongly on both axes. Still, a single benchmark’s cost-efficiency ranking is insufficient evidence
158 for selecting an agent on a new workload. Future agent leaderboards should use crossed model–
159 scaffold–benchmark designs with repeated runs and report both token- and dollar-based decision
160 frontiers.

161 References

- 162 Zhiyu Fan, Kirill Vasilevski, Dayi Lin, et al. SWE-effi: Re-evaluating software AI agent system
163 effectiveness under resource constraints. *arXiv preprint arXiv:2509.09853*, 2025. URL <https://arxiv.org/abs/2509.09853>.
164
- 165 Sayash Kapoor et al. AI agents that matter. *arXiv preprint arXiv:2407.01502*, 2024. URL <https://arxiv.org/abs/2407.01502>.
166
- 167 Sayash Kapoor et al. Holistic agent leaderboard: The missing infrastructure for AI agent evaluation.
168 *arXiv preprint arXiv:2510.11977*, 2026. URL <https://arxiv.org/abs/2510.11977>.
- 169 Franck Ndzomga. Efficient benchmarking of AI agents. *arXiv preprint arXiv:2603.23749*, 2026.
170 URL <https://arxiv.org/abs/2603.23749>.
- 171 Qi Qian, Chengsong Huang, et al. Benchmark²: Systematic evaluation of LLM benchmarks. *arXiv*
172 *preprint arXiv:2601.03986*, 2026. URL <https://arxiv.org/abs/2601.03986>.

A Frontier-label reproducibility audit

Table 3: Supplied HAL frontier labels compared with reconstructed frontiers by benchmark. Agreement is the fraction of displayed rows with identical Boolean labels.

Benchmark	Rows	Supplied	Nondominated	Hull	Supplied-ND agreement	Supplied-hull agreement
assistantbench	15	2	3	2	0.933	1.000
corebench_hard	45	5	10	5	0.889	1.000
gaia	32	4	6	4	0.938	1.000
online_mind2web	22	3	4	3	0.955	1.000
scicode	33	2	3	3	0.970	0.970
scienceagentbench	23	4	8	4	0.826	1.000
swebench_verified_mini	33	3	7	3	0.879	1.000
taubench_airline	26	3	3	4	1.000	0.962
usaco	13	4	5	4	0.923	1.000

B Reproducibility details

The analysis freezes the public CSV using the recorded source hash in the anonymous supplement. It uses exact public display labels, minimum pairwise overlap five, 5,000 configuration-resampling draws, and random seed 20260816. The released anonymous supplement should include the analysis code, input provenance manifest, generated tables, and commands needed to regenerate all figures.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction state that the analysis is descriptive, restricted to repeated public HAL display labels under a repeated displayed scaffold, and does not identify causal model or scaffold effects.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Section 6 discusses sparse coverage, display-label ambiguity, single-run evaluations, price dependence, missing token totals, limited domain comparisons, and frontier sensitivity.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[N/A\]](#)

Justification: The paper presents an empirical secondary analysis and no theoretical results.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: Sections 3–4 and the appendix specify the public source, frozen hash, inclusion criteria, matching unit, frontier definitions, overlap threshold, resampling procedure, and seed.

5. Open access to data and code

207 Question: Does the paper provide open access to the data and code, with sufficient instruc-
 208 tions to faithfully reproduce the main experimental results, as described in supplemental
 209 material?

210 Answer: [Yes]

211 Justification: An anonymized supplement will contain the analysis code, provenance mani-
 212 fest, and reproduction instructions; the underlying public data source is cited.

213 **6. Experimental setting/details**

214 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
 215 rameters, how they were chosen, type of optimizer) necessary to understand the results?

216 Answer: [Yes]

217 Justification: No models are trained. The data filtering, rank metrics, frontier definitions,
 218 resampling count, and seed are specified in Sections 3–4.

219 **7. Experiment statistical significance**

220 Question: Does the paper report error bars suitably and correctly defined or other appropriate
 221 information about the statistical significance of the experiments?

222 Answer: [Yes]

223 Justification: We report configuration-resampling sensitivity intervals and explicitly state
 224 that they are not rollout-level or population-model confidence intervals and are not used for
 225 hypothesis tests.

226 **8. Experiments compute resources**

227 Question: For each experiment, does the paper provide sufficient information on the com-
 228 puter resources (type of compute workers, memory, time of execution) needed to reproduce
 229 the experiments?

230 Answer: [N/A]

231 Justification: The work is a lightweight secondary analysis of public tabular data and uses
 232 no model training or agent rollouts.

233 **9. Code of ethics**

234 Question: Does the research conducted in the paper conform, in every respect, with the
 235 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

236 Answer: [Yes]

237 Justification: The work uses public aggregate evaluation data, attributes source authors,
 238 reports limitations, and makes no claims about people or protected groups.

239 **10. Broader impacts**

240 Question: Does the paper discuss both potential positive societal impacts and negative
 241 societal impacts of the work performed?

242 Answer: [N/A]

243 Justification: The work is a methodological analysis of public leaderboard measurements.
 244 Its practical implication is limited to more cautious agent selection and evaluation reporting.

245 **11. Safeguards**

246 Question: Does the paper describe safeguards that have been put in place for responsible
 247 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
 248 image generators, or scraped datasets)?

249 Answer: [N/A]

250 Justification: The paper releases no trained model, agent capability, or new high-risk dataset.

251 **12. Licenses for existing assets**

252 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
 253 the paper, properly credited and are the license and terms of use explicitly mentioned and
 254 properly respected?

255 Answer: [Yes]

256 Justification: The data and methodological sources are cited; the supplement records the
 257 upstream repository, commit, source path, and declared MIT license.

258 **13. New assets**

259 Question: Are new assets introduced in the paper well documented and is the documentation
 260 provided alongside the assets?

261 Answer: [Yes]

262 Justification: The analysis code and generated tables are documented by a source manifest,
 263 artifact manifest, tests, and reproduction commands.

264 **14. Crowdsourcing and research with human subjects**

265 Question: For crowdsourcing experiments and research with human subjects, does the paper
 266 include the full text of instructions given to participants and screenshots, if applicable, as
 267 well as details about compensation (if any)?

268 Answer: [N/A]

269 Justification: The study involves no crowdsourcing and no human subjects research.

270 **15. Institutional review board (IRB) approvals or equivalent for research with human**
 271 **subjects**

272 Question: Does the paper describe potential risks incurred by study participants, whether
 273 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
 274 approvals (or an equivalent approval/review based on the requirements of your country or
 275 institution) were obtained?

276 Answer: [N/A]

277 Justification: The study involves no human subjects research.

278 **16. Declaration of LLM usage**

279 Question: Does the paper describe the usage of LLMs if it is an important, original, or
 280 non-standard component of the core methods in this research?

281 Answer: [N/A]

282 Justification: The core method is deterministic analysis of public tabular data; no LLM is
 283 used as a component of the empirical method.