

# Does Cost-Efficiency Travel?

## Raw Cost Portability, Success-Adjusted Efficiency, and Matched Agent Trajectories

Anonymous Author(s)

August 19, 2026

### Abstract

Agent leaderboards increasingly report task success and dollar cost, inviting practitioners to select a “cost-efficient” system from a single benchmark. We ask whether this signal transfers across workloads, whether it reflects stable label-level cost propensity, and whether observable agent trajectories expose a scaffold contribution. First, using public Holistic Agent Leaderboard (HAL) data and a fixed displayed Generalist Agent scaffold, we compare repeated displayed model labels across six benchmarks. Raw dollar-cost ranks transfer in high-overlap pairs, but a pooled benchmark-plus-label log-cost model explains 82.0% of observed variation and residual cost ranks do not remain significant. Aggregate dollars per expected success are less portable than raw cost. Leave-one-benchmark-out prediction succeeds on four held-out workloads and fails on two lower-overlap workloads. Second, we analyze a separate public unencrypted matched trace sample: across 977 exact same-task/same-model pairs, OpenHands has lower median turns and tool calls than SWE-agent without a detected resolved-status difference. These studies establish that apparent agent cost efficiency mixes persistent displayed-label propensity, benchmark context, and scaffold trajectory structure. They do not identify tokens, provider prices, or dollar failure costs, which remain unavailable in the public snapshots.

## 1 Introduction

Agent evaluations increasingly expose a joint accuracy–cost tradeoff. This is a welcome correction to accuracy-only leaderboards: long-horizon systems can invoke tools, grow context, retry actions, and consume substantial resources. Yet a common practical inference remains underexamined: a label that is cheap on one benchmark is presumed to be an efficient deployment choice elsewhere. That inference conflates at least three mechanisms: persistent provider/model/configuration cost propensity, benchmark-specific demands, and scaffold execution behavior.

This paper studies the portability of observed agent cost structure. The primary study uses a frozen public HAL snapshot. Because public scaffold coverage is sparse, we do not claim a global model–scaffold–benchmark causal decomposition. Instead, we hold the displayed HAL Generalist Agent scaffold fixed and match exact public display labels across six workloads. We ask: (i) do raw dollar-cost ranks transfer; (ii) do they transfer beyond stable displayed-label cost propensity; (iii) does success adjustment change the answer; and (iv) is a cost ranking predictive on an unseen workload?

A second, explicitly separate trace study supplies a mechanism-oriented complement. The HAL public trace archives are encrypted, so they cannot answer how turns or tool calls differ. We therefore use a public unencrypted Open-SWE-Traces sample that permits exact task/model matching across two scaffolds. This trace extension measures structured trajectory burden proxies, not tokens or dollars.

Our contributions are:

1. a fixed-scaffold, coverage-aware analysis of raw cost, success-adjusted cost, and frontier portability across six HAL workloads;
2. held-out LOBO prediction and an observational displayed-label cost-propensity adjustment that separates persistent raw-cost structure from residual benchmark-specific structure;
3. a matched external trace extension showing substantial scaffold differences in turns and tool calls at fixed task and model; and
4. a reproducible artifact with source hashes, sensitivity analyses, and explicit boundaries on unavailable token, price, retry, and failure-cost data.

## 2 Related work

Kapoor et al. [2026] introduce HAL as cost-aware infrastructure for evaluating agents over multiple benchmarks. Kapoor et al. [2024] argue that agent evaluation should report accuracy and cost jointly, including a convex-hull interpretation of frontiers. Ndzomga [2026] study whether task subsets preserve ranks within agent benchmarks. SWE-Effi studies resource-constrained software-engineering systems and highlights scaffold–model interactions and expensive failures [Fan et al., 2025]. Benchmark<sup>2</sup> studies rank consistency across static LLM benchmarks [Qian et al., 2026]. Open-SWE-Traces releases large public SWE-agent and OpenHands trajectory corpora [Ahmad et al., 2026]. Our primary question differs: whether public observed agent cost rankings predict unseen workloads, and what they can and cannot mean as deployment-efficiency evidence.

## 3 Primary HAL study

### 3.1 Data and cohort

We use the public `all_leaderboards_costs_HAL.csv` distributed with Ndzomga [2026]. The frozen snapshot contains 242 displayed evaluations across nine benchmarks, 11 displayed scaffolds, and 29 displayed model labels. Most evaluations are single-run. Dollar cost and accuracy are available, but the CSV has no model/provider alias, prompt/output/cache token components, API-call count, tool-call count, retry count, or per-task trajectory records.

The only scaffold repeated broadly enough for the primary analysis is `HAL Generalist Agent`. We retain six workloads with adequate repeated-label overlap: `CORE-Bench Hard`, `GAIA`, `SciCode`, `ScienceAgentBench`, `SWE-bench Verified Mini`, and `TAU-bench Airline`. This yields 86 rows and 25 displayed labels. `USACO` has only one `Generalist` observation and is excluded. We match only exact public `Models` strings within the exact public `Agent Name`; this is not a controlled base-model identity claim.

### 3.2 Metrics

For every benchmark pair with at least five shared labels, we compare ranks using Spearman  $\rho$  and Kendall  $\tau_b$ . We report overlap  $N$ . Pairwise rank intervals use 10,000 configuration-resampling draws, interpreted as finite-label sensitivity intervals rather than rollout uncertainty. Holm–Bonferroni correction is applied separately to 15 accuracy and 15 raw-cost tests.

We define aggregate dollars per expected success as

$$\text{CPS}_{i,b}(\epsilon) = \frac{C_{i,b}}{\max(A_{i,b}, \epsilon)},$$

Table 1: High-overlap raw dollar-cost rank transfer and residual cost-rank sensitivity. Residual ranks remove observed displayed-label and benchmark log-cost effects; they are not list-price-adjusted costs.

| Pair          | N  | Raw $\rho$ | Residual $\rho$ | Perm. $p$ |
|---------------|----|------------|-----------------|-----------|
| CORE-GAIA     | 14 | 0.93       | 0.23            | 0.426     |
| CORE-SWE-mini | 14 | 0.74       | -0.49           | 0.079     |
| CORE-TAU      | 12 | 0.78       | -0.03           | 0.921     |
| GAIA-SWE-mini | 17 | 0.89       | 0.03            | 0.917     |
| GAIA-TAU      | 14 | 0.86       | 0.13            | 0.648     |
| SWE-mini-TAU  | 14 | 0.82       | -0.16           | 0.570     |

where  $C_{i,b}$  is displayed aggregate dollar cost and  $A_{i,b} \in [0, 1]$  is accuracy. The primary floor is  $\epsilon = 0.01$ ;  $\epsilon = 0.001$ ,  $\epsilon = 0.05$ , and zero-accuracy exclusion are reported as sensitivities. CPS is an aggregate ratio, not mean spending among successful trajectories.

We compute standard discrete nondominated frontiers and a HAL-inspired origin-anchored convex-hull reconstruction. A 5% tolerance sensitivity requires both 5% lower cost and 5% higher accuracy for dominance. It is a near-tie stress test, not a measurement-error model.

### 3.3 Cost-propensity and held-out models

To assess whether raw cost transfer is dominated by stable observed label propensity, we fit

$$\log C_{i,b} = \mu + \alpha_i + \beta_b + \epsilon_{i,b},$$

where  $\alpha_i$  is a displayed-label fixed effect and  $\beta_b$  is a benchmark effect. This is an observational adjustment, not a provider-list-price adjustment. Stable label propensity can include prices, reasoning settings, routing, token behavior, cache treatment, and unrecorded setup.

For leave-one-benchmark-out (LOBO) prediction, we hold out a benchmark, center log cost within each training benchmark, average each label’s centered cost over the remaining benchmarks, and predict the held-out rank. Labels must occur in at least two training workloads. We assess each held-out rank correlation with a 20,000-draw within-benchmark label-pairing permutation null.

## 4 Primary results

### 4.1 Raw dollar-cost rank transfers, but residual cost rank does not

Across the six pairs with at least 12 shared labels, raw dollar-cost Spearman correlation ranges from 0.74 to 0.93, with mean 0.84. These associations survive the pre-specified raw-cost multiplicity and permutation checks. However, the pooled cost-propensity model explains 82.0% of log-cost variation (adjusted  $R^2 = 72.7\%$ ). After removing benchmark and displayed-label cost propensity, no high-overlap residual-cost rank association is significant under a 20,000-draw permutation test. The result is not that cost is irrelevant; it is that raw cost portability largely reflects a stable observed label-level component rather than demonstrated transferable trajectory behavior.

### 4.2 Success adjustment is less portable

Across the same six high-overlap pairs, mean raw-cost rank correlation is 0.84 whereas mean CPS rank correlation under the 1% floor is 0.59. CPS portability declines in five pairs. GAIA-SWE-mini changes from 0.89 for raw cost to 0.29 for CPS; CORE-SWE-mini changes from 0.74 to 0.32; SWE-mini-TAU changes




Figure 1: Raw-cost versus CPS rank transfer for high-overlap pairs. CPS uses the primary 1% denominator floor.

from 0.82 to 0.45. CORE-TAU is an exception, with CPS 0.83 versus raw cost 0.78. Thus a label that is predictably cheap is not necessarily predictably cheap per expected success on a different workload.

### 4.3 LOBO identifies portable and nonportable raw-cost settings

Stable observed label propensity predicts held-out raw cost ranks for CORE-Bench Hard ( $\rho = 0.78$ ,  $N = 14$ ,  $p = 0.0016$ ), GAIA ( $\rho = 0.87$ ,  $N = 16$ ,  $p < 0.0001$ ), SWE-mini ( $\rho = 0.90$ ,  $N = 16$ ,  $p < 0.0001$ ), and TAU Airline ( $\rho = 0.83$ ,  $N = 14$ ,  $p = 0.0004$ ). It does not predict SciCode ( $\rho = -0.17$ ,  $N = 9$ ,  $p = 0.6801$ ) or ScienceAgentBench ( $\rho = -0.64$ ,  $N = 7$ ,  $p = 0.1376$ ). The latter estimates are sparse; they establish benchmark-specific heterogeneity rather than a scientific-domain causal explanation.

### 4.4 Frontier membership is non-universal

No label tested on at least five workloads appears on every observed accuracy–cost frontier under discrete nondominance, the HAL-inspired hull, or 5% tolerance. The stronger claim that no label appears on more than 4/6 frontiers is not robust: under tolerance, o4-mini Low appears on 5/6. We therefore retain only the non-universal-winner conclusion.

## 5 External matched trajectory study

### 5.1 Data and design

The HAL aggregate CSV cannot identify execution mechanisms because its traces are encrypted. We therefore conduct an external, separate matched trace analysis using public unencrypted Open-SWE-Traces data. We select a documented four-shard boundary sample for MiniMax-M2.5 under OpenHands and SWE-agent. The boundary sample is reproducible but not a random sample of the full corpus. We exact-match task ID and model, exclude duplicate task/model/scaffold records, and obtain 977 pairs.




Figure 2: LOBO predicted versus observed raw-cost ranks. Blue bars differ from the 20,000-draw label-pairing permutation null at 0.05; orange bars do not.

The trace schema contains resolved status, message roles, explicit tool calls, message text, and reasoning text. It does not contain billed tokens, dollars, provider usage, retries, or latency. We measure trajectory turns, assistant turns, explicit tool calls, tool-result turns, total message characters, and reasoning-message characters. Character counts are tokenizer-free text-volume proxies, not tokens or costs. Paired comparisons use two-sided Wilcoxon signed-rank tests.

## 5.2 Scaffold changes trajectory burden at fixed task and model

OpenHands has lower medians across the matched sample: 123 versus 149 trajectory turns, 61 versus 74 assistant turns, 61 versus 74 tool calls, and 176,708 versus 194,263 total trajectory characters. The corresponding paired Wilcoxon tests are significant (turns and tool calls  $p = 1.9 \times 10^{-49}$ ; characters  $p = 2.0 \times 10^{-12}$ ). Resolved status has no detected paired difference ( $p = 0.197$ ). This evidence shows that, at fixed task and model in this external sample, scaffold choice can materially alter observable trajectory structure without a detected success difference. It does not claim lower token use, lower dollar cost, or a causal explanation of HAL cost rankings.

## 6 Discussion

The combined results clarify what a cost-aware agent leaderboard can and cannot establish. Raw dollar-cost rankings are predictive for some held-out workloads, but their portability is largely associated with stable observed displayed-label cost propensity. CPS portability is weaker, so raw cheapness is not interchangeable with deployment efficiency. Independent matched traces further show that scaffolds change measurable execution structure, even when task and model are fixed.

This is a secondary-data analysis with clear construct boundaries. The primary CSV has sparse, non-random coverage, mostly single-run rows, and public display labels that omit prompts, tools, budgets, and routing. Cost-propensity adjustment is not provider-price control. CPS is not successful-run cost. The external trace study is not merged with HAL, uses a nonrandom boundary sample, and uses burden proxies




Figure 3: Matched MiniMax-M2.5 task pairs in the Open-SWE boundary sample ( $N = 977$ ). Markers show scaffold medians on a log scale for trace proxies, not tokens or dollars.

rather than billing fields. We cannot compute failure-cost ratios, tail dollar costs, tokens, cache usage, retries, or a true three-dimensional accuracy–cost–reliability frontier from public data. These limitations motivate public release of aggregate usage and outcome telemetry with leaderboard results.

## 7 Conclusion

Observed dollar-cost portability is real in several high-overlap agent benchmark pairs, but it should not be read as a universal deployment-efficiency signal. It is partly persistent label-level cost propensity, weakens after success adjustment, and coexists with substantial scaffold differences in matched trajectory structure. Agent leaderboards should report not only accuracy and raw dollar cost, but also success-adjusted outcomes and trace-level usage fields needed to distinguish price, scaffold, and execution mechanisms.

## References

- Wasi Uddin Ahmad, Nikolai Ludwig, Somshubra Majumdar, and Boris Ginsburg. Open-SWE-traces: Advancing dual-mode multilingual distillation for software engineering agents. *arXiv preprint arXiv:2606.16038*, 2026. URL <https://arxiv.org/abs/2606.16038>.
- Zhiyu Fan, Kirill Vasilevski, Dayi Lin, et al. SWE-effi: Re-evaluating software AI agent system effectiveness under resource constraints. *arXiv preprint arXiv:2509.09853*, 2025. URL <https://arxiv.org/abs/2509.09853>.
- Sayash Kapoor et al. AI agents that matter. *arXiv preprint arXiv:2407.01502*, 2024. URL <https://arxiv.org/abs/2407.01502>.
- Sayash Kapoor et al. Holistic agent leaderboard: The missing infrastructure for AI agent evaluation. *arXiv preprint arXiv:2510.11977*, 2026. URL <https://arxiv.org/abs/2510.11977>.

Franck Ndzomga. Efficient benchmarking of AI agents. *arXiv preprint arXiv:2603.23749*, 2026. URL <https://arxiv.org/abs/2603.23749>.

Qi Qian, Chengsong Huang, et al. Benchmark<sup>2</sup>: Systematic evaluation of LLM benchmarks. *arXiv preprint arXiv:2601.03986*, 2026. URL <https://arxiv.org/abs/2601.03986>.