
Observability Debt in Agent Evaluation: Toward a Telemetry Layer for AgenticOS

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agentic AI systems are evaluated today through three disjoint public
2 infrastructures—leaderboards, trace corpora, and telemetry logs—none of the pub-
3 lic sources we examine jointly exposes the fields a resource-aware operating layer
4 would need: task identity, model identity, scaffold and policy-configuration iden-
5 tity, outcome, token/cache usage, retries, latency, and billing. We measure its
6 consequences across three studies on public data. First, using the Holistic Agent
7 Leaderboard (HAL), raw dollar-cost rankings transfer across some benchmark
8 pairs (Spearman $\rho=0.74\text{--}0.93$) but are mostly absorbed by a workload-persistent
9 displayed-label cost propensity (pooled model $R^2=82.0\%$); success-adjusted cost
10 is markedly less portable (mean $\rho=0.59$ vs. 0.84). Second, using 977 exact task-
11 model-matched trajectory pairs from Open-SWE-Traces, scaffold choice changes
12 observable execution burden at fixed task and model (OpenHands/SWE-agent turn
13 ratio 0.85 , $[0.82, 0.87]$), and the widely assumed “expensive failure” premium is
14 not detected under exact task matching (80%-power threshold: premium fractions
15 ≥ 0.65 ; observed $0.44\text{--}0.49$). Third, auditing two independent telemetry corpora
16 (8,058 and 1,017 sessions), we characterize the resource anatomy—93.7% mean
17 cache-read share, $4.29\times$ conditional tool-error cascades replicated across sources,
18 top-1%-of-sessions holding 48.3% of input-token volume—that a scheduler or
19 router would need but that no benchmark currently attaches to an outcome. These
20 studies are complementary rather than linked: Studies 1–2 show that transferred
21 raw-cost rankings and unconditioned failure premiums confound *which workloads*
22 *were attempted with how agents behaved on them*, while Study 3 shows that the ex-
23 ecution variables hidden beneath those aggregates are themselves heterogeneous,
24 heavy-tailed, and state-dependent. From this evidence we derive a joint telemetry
25 schema spanning identity, execution, resources, control, and outcome, and argue
26 it belongs at the operating-system layer: controlling execution and allocating re-
27 sources requires observing what those allocations depend on.

1 Why This Belongs at the OS Layer

29 Agentic AI systems persist state, invoke tools, retry failed actions, and accumulate context across
30 long horizons. Reasoning about their resource behavior—for scheduling, routing, cost budgeting,
31 or governance—requires the same unified telemetry that operating systems have long provided for
32 conventional workloads: a common schema linking *what ran, on what, under which policy, with*
33 *what outcome, and at what resource cost*. We use *observability debt* to name the accumulated
34 inability of public agent-evaluation infrastructure to jointly attribute resource consumption to task,
35 model, execution policy, outcome, and system-level events. In every public source we examined,
36 this debt prevents reliable attribution of observed resource differences to workload, model, scaffold,
37 configuration, or execution behavior.




Figure 1: **The agentic OS control loop.** Each stage consumes exactly the schema fields derived in Section 6; the feedback arrow hosts scheduling, routing, budgeting, and halting decisions that today run on aggregate dollars because no examined public source closes this loop (Table 2).

38 What makes this an *operating-system* responsibility rather than better benchmark bookkeeping? An
39 agentic OS occupies a specific functional position. It observes *identity* (which task, model, scaf-
40 fold, and policy configuration is running), *execution* (which tools act, what context accumulates),
41 *resources* (tokens, cache, time, memory, money), *control* (retry, terminate, route, budget), and *out-*
42 *come* (did the task succeed). The same resource-accounting principle motivates an OS-level treat-
43 ment: the layer that controls execution and allocates resources needs access to the telemetry required
44 to make those decisions. Without per-round token components, a memory manager cannot budget
45 context growth; without configuration identity, a router cannot learn *why* one label is cheap; without
46 error-status streams, a governor cannot halt a spiraling session. Today, each of these decisions is
47 made implicitly from aggregate dollar figures (Figure 1).

48 We treat three public data sources as stand-ins for what this observability layer would unify: a
49 cost-and-outcome leaderboard (HAL), a task/model/scaffold-matched trajectory corpus (Open-SWE-
50 Traces), and raw session-level usage telemetry (two independent corpora). Each exposes a different
51 partial view. HAL links dollar cost to outcome but encrypts execution traces. Open-SWE-Traces ex-
52 poses exact task/model/scaffold matching and resolved status but no tokens or billing. The telemetry
53 corpora expose token, cache, and tool-call dynamics but no outcome or billing. **No public source we**
54 **examined jointly exposes task, model, scaffold and policy configuration, outcome, token/cache**
55 **usage, retries, latency, and billing**, the joint schema a resource-aware OS layer needs (Table 2).
56 Importantly, combining these sources does not create record-level observability: their different ex-
57 ecution units, identifiers, and release boundaries prevent resource use, outcome, and billing from
58 being joined into a single per-execution record.

59 Our three studies are therefore complementary rather than linked—each probes a different dimen-
60 sion of the gap, and none can be joined to another at the record level. (1) A coverage-aware cost-
61 portability study (HAL) shows that aggregate dollar cost conflates displayed-label propensity with
62 benchmark difficulty and decorrelates under success adjustment. (2) A matched-trajectory study
63 shows that execution structure differs measurably across scaffolds, while a popular failure-cost
64 heuristic turns out to measure workload composition. (3) A cross-source telemetry audit shows
65 that the execution variables hidden beneath aggregate cost are themselves highly heterogeneous,
66 heavy-tailed, and state-dependent. Together they triangulate one gap; no pair of them jointly identi-
67 fies a causal mechanism. This paper is an empirical requirements study for an AgenticOS telemetry
68 layer, rather than an implementation of the layer itself.

69 2 Related Work

70 Kapoor et al. (2026) introduce HAL as cost-aware evaluation infrastructure spanning coding, web,
71 science, and customer-service tasks (21,730 rollouts, roughly \$40,000 in evaluation cost). Kapoor et
72 al. (2024) motivate joint accuracy–cost reporting and convex-hull frontier analysis, which we adopt
73 as one of three decision-frontier definitions. Ndzomga (2026) asks whether task subsets preserve
74 agent rankings within a single benchmark; we instead ask whether an *observed cost signal* pre-
75 dicta a *different* workload. Benchmark² studies cross-benchmark rank consistency for static LLM
76 benchmarks; our setting adds interactive scaffolds, dollar cost, success adjustment, and a held-out
77 workload test. SWE-Effi reports that failed attempts consume substantially more resources than suc-
78 cessful ones; our matched design tests that claim at fixed task and does not detect the premium once
79 task-difficulty confounding is removed (Section 4). Open-SWE-Traces supplies the exact matching
80 that makes this test possible but lacks token/billing fields, motivating Section 6. Commercial trac-
81 ers expose per-deployment spans but are not cross-benchmark evaluation, and public releases strip
82 execution detail, so the per-request visibility our audited sources lack is exactly what we quantify.




Figure 2: **Raw cost transfers; success-adjusted cost does not.** Pairwise Spearman ρ for benchmark pairs sharing ≥ 12 labels. Raw dollar cost is portable (mean 0.84); cost per expected success is not (mean 0.59), only three CPS pairs survive correction.

3 Study 1: Does a Cost Signal Travel Across Workloads?

3.1 Data and cohort

We use the public HAL CSV snapshot (Ndzomga, 2026): 242 displayed evaluations across nine benchmarks, 11 scaffolds, 29 labels, supplying accuracy and aggregate dollar cost but no token components, retries, or per-task trajectories—the gap this paper is about.

The only broadly repeated scaffold is the displayed HAL Generalist Agent; fixing it and matching exact label strings yields six adequately overlapping workloads (CORE-Bench Hard, GAIA, SciCode, ScienceAgentBench, SWE-bench Verified Mini, TAU-bench Airline), a cohort of 86 rows and 25 labels.

3.2 Raw cost transfers; the mechanism behind it does not

Across six pairs with ≥ 12 shared labels, raw dollar-cost Spearman ρ ranges 0.74–0.93 (mean 0.84), surviving Holm-corrected permutation testing (Figure 2). A pooled label-plus-benchmark log-cost fixed-effects model explains $R^2=82.0\%$ (adjusted 72.7%); once label and benchmark propensity are removed, **no high-overlap residual-cost association survives** a 20,000-draw permutation test.

Leave-one-benchmark-out (LOBO) prediction succeeds for CORE ($\rho=0.78$, $N=14$, $p=0.0016$), GAIA (0.87, $N=16$), SWE-mini (0.90, $N=16$), TAU (0.83, $N=14$). It fails for SciCode (unstable, $[-0.50, 0.14]$) and ScienceAgentBench (directionally negative, $[-0.77, -0.43]$).

Interpretation: the portable component behaves as a workload-persistent displayed-label propensity—plausibly reflecting provider pricing, reasoning-mode defaults, or routing policy—that a router could in principle learn and reuse. but public data cannot identify *which*, since the distinguishing fields are not exposed.

3.3 Success-adjusted cost is a different, less portable ranking

Cost per expected success, $CPS_{i,b}(\epsilon) = C_{i,b} / \max(A_{i,b}, \epsilon)$ with primary floor $\epsilon=0.01$ (sensitivities at 0.001, 0.05, and zero-exclusion in the supplement), has mean $\rho=0.59$. Only three pairs survive correction: CORE–GAIA (0.87, $[0.63, 0.96]$), CORE–TAU (0.83, $[0.31, 1.00]$), GAIA–TAU (0.78, $[0.42, 0.94]$). The clearest raw-vs-CPS gap is GAIA–SWE-mini (0.60, $[0.14, 1.13]$): a label cheap in dollars is not reliably cheap per expected success elsewhere. No label tested on ≥ 5 workloads sits on every accuracy–cost frontier under discrete nondominance, convex hull, or 5% tolerance (o4-mini Low reaches 5/6 under tolerance; no universal winner under any strict definition).

Implication: Routing on displayed dollar cost therefore optimizes a signal whose portability degrades substantially once success is taken into account.

4 Study 2: What Matched Trajectories Reveal That Aggregates Cannot

HAL’s own trace archives are encrypted, so Study 1 cannot say *why* costs differ. We use an external unencrypted corpus—977 exact task–model pairs from Open-SWE-Traces for MiniMax-M2.5 under two scaffolds (OpenHands, SWE-agent), a reproducible but nonrandom four-shard boundary sample.




Figure 3: (a) Scaffold burden on 977 matched pairs: OpenHands uses $\sim 15\%$ fewer median turns/tool calls than SWE-agent at fixed task and model. (b) The failure premium is not detected under exact task matching (80% power against $\geq 15pp$). Left: naive cross-task ratios reproduce folklore (19–36% heavier failures). Right: among 107 discordant pairs, failed side heavier only 47–52/107.

Table 1: **Telemetry anatomy across two independent corpora.** Cache metrics are within-corpus; a share and a total-volume ratio are not interconvertible.

	TraceLab	MIMO Claude Code
Sessions / LLM rounds	8,058 / 665,453	1,017 / 4,690
Median input / output tokens per round	132,092 / 249	n/a (constant-tuple audit)
Cache-read vs. uncached input	93.7% mean share ($\approx 14.9\times$)	753.0M / 46.3M tokens = $16.3\times$
Median final-vs-initial context	$3.6\times$ (P90 $9.9\times$)	n/a
Top-1% / top-10% input-token share	48.3% / 88.2%	— / 37.9%
Error cascade $P(\text{err} \text{err})/P(\text{err} \text{succ})$	$21.0/4.9\% = 4.29\times$	4.3% vs. 1.0% ($4.33\times$)
Errored calls retried same-tool / retry succeeds	88.3% / 77.9%	—
Consecutive-error run length (median/P90/max)	1 / 3 / 368	—
Tool-call latency, success vs. error (median)	140 ms / 313 ms	—
Tool-result size, success vs. error (median chars)	602 / 208	—

118 The schema exposes resolved status, roles, tool calls, and text but not billed tokens or dollars; we
 119 treat turns, tool calls, and characters as trace-burden proxies.

120 4.1 Scaffold choice changes observable burden at fixed task

121 OpenHands has lower matched medians: 123 vs. 149 trajectory turns (ratio 0.85, $[0.82, 0.87]$), 61 vs.
 122 74 tool calls (0.85, $[0.82, 0.86]$), 176,708 vs. 194,263 characters (0.95, $[0.92, 0.98]$). Resolved status
 123 shows no detected difference ($p=0.197$): scaffolds differ in execution structure without a detected
 124 accuracy trade-off here.

125 4.2 The “expensive failure” premium is not detected under exact task matching

126 Comparing failed vs. resolved trajectories *across* tasks reproduces SWE-Effi’s pattern (19–36%
 127 heavier failures). But among the **107 discordant pairs**—identical task and model, one scaffold
 128 resolved and the other failed—the failed side is heavier only 47–52 of 107 on every metric (bi-
 129 nomial $p \geq 0.25$). A pre-specified power calculation gives 80% power against a premium fraction
 130 ≥ 0.65 ; the observed 0.44–0.49 does not resolve smaller effects ($< 15pp$).

131 **Interpretation:** hard tasks make agents both burn more *and* fail more. An unconditioned failure
 132 multiplier measures *which tasks were attempted*, not a property of failing runs. Task-conditioned
 133 accounting requires exactly the joint schema no single public source provides.

134 5 Study 3: The Telemetry an OS Layer Would Need to See Directly

135 Studies 1–2 show aggregate cost hides *why* an agent is expensive and that execution structure varies
 136 across systems at fixed task and model. This section audits raw execution telemetry across two
 137 independent coding-agent corpora (Table 1). **TraceLab:** 8,058 sessions, 665,453 rounds, 743,819
 138 tool calls. **MIMO Claude Code:** 1,017 sessions, 4,690 rounds. Neither links usage to outcome or
 139 billing.

140 **Context is the dominant resource, not generation.** The median TraceLab round carries 132,092
 141 input tokens against 249 output tokens (93.7% mean cache-read share); median final-round context




Figure 4: **Heavy-tailed burden.** Top-1%/top-10% summed-input concentration and median session input by observed-error bin (0.34M at zero to 14.65M at 6+).

is $3.6\times$ initial (P90 $9.9\times$). A dollar total cannot separate price- from cache-policy- or accumulation-driven spend. Tool verbosity drives it: median next-input growth is 443 / 1,412 / 5,420 tokens below / within / above 10k tool-result characters.

Resource burden is heavy-tailed. The top 1% of sessions hold 48.3% of summed input volume; top 10% hold 88.2%. Median session input rises from 0.34M tokens at zero observed tool errors to 14.65M at six or more; mean-based budgeting would miss this tail-driven concentration (Figure 4).

Error cascades replicate across independent sources. After a successful call the next errs 4.9%, after an error 21.0% ($4.29\times$); MIMO replicates at $4.33\times$ (4.3% vs. 1.0%). Same-tool retry is common (88.3% get one, 77.9% succeed) but the tail is extreme (consecutive-error runs: median 1, P90 3, max 368). Failures are also slower yet less informative: 313 ms vs. 140 ms (P95 85.6 s vs. 15.8 s), returning 208 vs. 602 characters.

Data-quality boundary. MIMO repeats one constant per-round usage tuple, so its within-session token dynamics are not validated; we use it only for session counts, totals, concentration, and content-derived error sequences.

6 An Evidence-Derived Joint Telemetry Schema

Table 2 shows the coverage gap; the missing rows are exactly those Sections 3–5 show to be load-bearing. Each field is tied to an analysis above or to the OS control function it unblocks (*evidence-derived*: usefulness and necessity for those functions, not formal minimality). Keyed to that function:

1. (*Attribution / exact matching*) **Task and model identity** (exact IDs, not display strings)—enables exact-matched designs; without it the Section 4 test is infeasible on public data.
2. (*Routing between scaffolds*) **Scaffold/harness identity and version**—attributes burden differences to scaffold design rather than model behavior; input to routing.
3. (*Configuration attribution*) **Configuration and policy identity** (reasoning mode, sampling params, budgets, routing/cache policy)—distinct from model identity; Study 1’s unexplained propensity component is plausibly the shadow of these missing fields.
4. (*Success-conditioned accounting*) **Outcome/resolved status**—conditions any cost statistic on success, closing the CPS portability gap.
5. (*Memory / budget management*) **Token components per round (input, output, cache-read, cache-creation)**—separates price from accumulation and cache policy; lets a memory layer budget context growth.
6. (*Governance / halting*) **Tool-call counts, retries, error status per call**—detects and prices $4.29\times$ cascade dynamics; a design motivation for governance policies that halt spiraling sessions—motivated here, not experimentally validated.
7. (*Real-time scheduling*) **Latency per call and per session**—present only as unlinked tool-call durations in one corpus; required for real-time scheduling.
8. (*Procurement / auditing*) **Billing, linked to all of the above**—the one field HAL exposes well, and exactly the field that is uninterpretable without the others.

Figure 1 places these fields at the stages of the control loop they close; Table 2 shows why, on public data alone, the loop does not close.

No public source we examined exposes this joint schema. We do not propose a storage format or API—a systems-design question outside our scope—but the measured consequences of its absence argue it belongs at a shared operating-system layer rather than being re-derived piecemeal inside each benchmark, where today it is not derived at all.

Table 2: **The observability gap.** Coverage per public source examined: ■ exposed, □ partial (display-label-only, degraded fidelity, or inferable), – absent. No source covers the joint schema; the union leaves latency unlinked and round-level token fidelity unverifiable outside one corpus.

Field	HAL	Open-SWE	TraceLab	MIMO
Task identity	□	■	–	–
Model identity	□	■	□	■
Scaffold identity/version	■	■	□	■
Configuration / policy identity	□	–	–	–
Outcome / resolved status	■	■	–	–
Tokens / cache components	–	–	■	□
Retries / error status	–	□	■	□
Latency	–	–	■	–
Billing	■	–	–	–

7 Limits

This is a secondary-data analysis with clear boundaries. Display labels do not fully identify prompts, tools, budgets, or harness versions; the HAL cohort is sparse, unbalanced, mostly single-run. Propensity adjustment is observational, not list-price control; CPS is an aggregate ratio, not mean cost of successful runs. Study 2 uses one model and a nonrandom sample whose proxies are not tokens or dollars, powered only against strong premiums ($\geq 15\text{pp}$). Study 3 is observational and unlinked to outcomes or billing; MIMO’s per-round fields are unverifiable and used only where supported, and its design implications are motivations, not evaluated policies. Each limit is itself an instance of the missing-field problem.

8 Conclusion

Across three studies, aggregate resource signals cannot be reliably attributed to workload, execution, or outcome without telemetry fields the public sources we examined do not jointly expose. Raw-cost rankings transfer across workloads but are strongly associated with a workload-persistent displayed-label cost propensity; success-adjusted cost reorders the ranking (three of six pairs survive correction); the assumed failure-cost premium is not detected under exact task matching. No examined source jointly links task, model, scaffold and configuration, outcome, token/cache usage, retries, latency, and billing. The layer that controls execution and allocates resources is the natural owner of this telemetry, so we argue the schema belongs at a shared operating-system layer for agentic AI rather than being reconstructed piecemeal inside each leaderboard or trace release.

References

- S. Kapoor et al. AI agents that matter. *arXiv preprint arXiv:2407.01502*, 2024.
- W. U. Ahmad, N. Ludwig, S. Majumdar, and B. Ginsburg. Open-SWE-traces: Advancing dual-modal multilingual distillation for software engineering agents. *arXiv preprint arXiv:2606.16038*, 2026.
- Z. Fan, K. Vasilevski, D. Lin, et al. SWE-eff: Re-evaluating software AI agent system effectiveness under resource constraints. *arXiv preprint arXiv:2509.09853*, 2025.
- S. Kapoor et al. Holistic agent leaderboard: The missing infrastructure for AI agent evaluation. *arXiv preprint arXiv:2510.11977*, 2026.
- Franck Ndzomga. Efficient benchmarking of AI agents. *arXiv preprint arXiv:2603.23749*, 2026.
- Q. Qian, C. Huang, et al. Benchmark²: Systematic evaluation of LLM benchmarks. *arXiv preprint arXiv:2601.03986*, 2026.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: Claims in the abstract and introduction are supported by the three public-data studies (Sections 3–5) via permutation/Holm tests, bootstrap and binomial intervals, and pre-specified power analyses (Sections 3–4); limitations are stated in Section 7 (Limits).

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 7 (Limits) discusses the single-model matched sample, nonrandom trace sample, proxy (non-dollar) measures, observational telemetry, and the unverifiable MIMO per-round fields.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: Empirical paper; decompositions are observational fixed-effects models with stated assumptions (Section 3). No theoretical results requiring proof.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All input datasets are public (HAL CSV, Open-SWE-Traces, TraceLab, MIMO) and are cited; analysis code and frozen outputs will be released in anonymized form upon acceptance, with cohort construction, resampling, and permutation procedures described in Sections 3–4.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All input datasets are public research releases and are cited in the References; analysis code and frozen outputs will be released in anonymized form upon acceptance, with sufficient instructions to reproduce the main results.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Settings stated: CPS floor $\epsilon=0.01$ (sensitivities 0.001, 0.05, zero-exclusion); 20,000-draw permutations; bootstrap seeds; minimum overlap $n \geq 12$ (Sections 3–4).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: We report permutation and bootstrap intervals, Holm-corrected permutation tests, and binomial sign tests with a pre-specified 80% power analysis (Sections 3–4).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Lightweight computations over public CSV/trace files on a standard workstation (Apple M-series, 32 GB RAM); minutes of runtime, no GPU or large-scale compute.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Secondary analysis of public evaluation artifacts; no human subjects, harmful content, or privacy-sensitive data involved.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 7 notes improved evaluation infrastructure is positive (resource awareness, cost/efficiency) with no direct negative impact; documented gaps are missing-field problems, not enabled harms.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No high-risk assets are released; we release analysis code and a proposed telemetry schema, not models or scraped datasets with misuse potential.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: HAL CSV, Open-SWE-Traces, TraceLab, and MIMO are public research releases and are cited in the References.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: We release the analysis pipeline, figures, and proposed telemetry schema (Sections 5–6), documented in the accompanying repository.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: No crowdsourcing or research with human subjects is involved.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

538 Justification: No human subjects are involved.
539 Guidelines:
540 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
541 with human subjects.
542 • Depending on the country in which research is conducted, IRB approval (or equiva-
543 lent) may be required for any human subjects research. If you obtained IRB approval,
544 you should clearly state this in the paper.
545 • We recognize that the procedures for this may vary significantly between institutions
546 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
547 guidelines for their institution.
548 • For initial submissions, do not include any information that would break anonymity
549 (if applicable), such as the institution conducting the review.

550 16. Declaration of LLM usage

551 Question: Does the paper describe the usage of LLMs if it is an important, original, or
552 non-standard component of the core methods in this research? Note that if the LLM is used
553 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
554 scientific rigor, or originality of the research, declaration is not required.

555 Answer: [Yes]

556 Justification: LLM assistants aided drafting/editing; all empirical claims and numbers were
557 verified by the authors against code and data, without impacting core methodology or orig-
558 inality of the research.

559 Guidelines:

- 560 • The answer [N/A] means that the core method development in this research does not
561 involve LLMs as any important, original, or non-standard components.
- 562 • Please refer to our LLM policy in the NeurIPS handbook for what should or should
563 not be described.