
Does Cost-Efficiency Travel? Rank Transfer and Frontier Stability in Fixed-Scaffold Agent Evaluation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agent leaderboards increasingly report task success and dollar cost, inviting practitioners to
2 select a “cost-efficient” system from a single benchmark. We test whether this signal transfers
3 across workloads and whether any transfer reflects stable model-label cost propensity. Using
4 public Holistic Agent Leaderboard (HAL) results, we hold the displayed HAL Generalist
5 Agent scaffold fixed and compare repeated displayed model labels across six benchmarks.
6 Raw dollar-cost ranks transfer robustly in several high-overlap pairs: among pairs with
7 at least 12 shared labels, mean pairwise Spearman correlation is 0.84. A pooled log-cost
8 model with benchmark and displayed-label fixed effects explains 82 percent of observed
9 log-cost variance; after removing this stable observed label propensity, no high-overlap
10 residual-cost rank association is significant under a 20,000-draw permutation test. Aggregate
11 CostPerSuccess, dollar cost divided by $\max(\text{accuracy}, 1 \text{ percent})$, has lower mean pairwise
12 rank correlation, 0.59. Finally, an external matched unencrypted trace sample of 977 same-
13 task/same-model pairs shows that OpenHands uses fewer median turns and tool calls than
14 SWE-agent without a detected resolved-status difference. Thus dollar cost mixes stable
15 label propensity and scaffold trajectory structure; raw cost transfer alone is not portable
16 deployment efficiency. No broad label lies on every cost–accuracy frontier under three
17 definitions.

18 1 Introduction

19 Agent benchmarks evaluate systems that reason over multiple steps, invoke tools, and interact with environments.
20 Reported behavior depends on the language model, scaffold, tools, prompts, execution limits, and task environ-
21 ment. Evaluation is also expensive: HAL reports 21,730 rollouts across nine benchmarks and roughly \$40,000 in
22 cost [Kapoor et al., 2026]. Cost-aware leaderboards therefore plot accuracy against dollar cost and recommend
23 cost–accuracy frontiers.

24 A practical question follows: if a displayed model label is efficient on one agent benchmark, is it a promising
25 low-cost choice elsewhere? This may hold when capability and resource use are portable, but can fail when
26 tool calls, retries, or stopping behavior differ by environment. We study this question descriptively, holding the
27 publicly displayed HAL Generalist Agent scaffold fixed and matching exact public model-label strings across
28 workloads.

29 Our design is coverage-aware. Most HAL scaffolds occur on one benchmark, making a global model–scaffold
30 decomposition unsupported. The Generalist Agent is the only broadly repeated scaffold. We retain six bench-
31 marks with adequate pairwise overlap and analyze accuracy rank, dollar-cost rank, aggregate CostPerSuccess
32 rank, and frontier membership only among shared labels. CostPerSuccess is defined as $C / \max(A, 0.01)$, where
33 A is accuracy as a fraction; the supplement reports a 5

34 2 Related work

35 Kapoor et al. [2026] introduce HAL as a cost-aware leaderboard across coding, web, science, and customer-
36 service tasks. Kapoor et al. [2024] argue for joint accuracy–cost evaluation of agents. SWE-Effi studies resource-
37 constrained software-engineering agents and reports scaffold–model interaction within one domain [Fan et al.,

Table 1: Higher-overlap pairs. Brackets are configuration-resampling sensitivity intervals; ND is nondominated-frontier Jaccard similarity on the shared-label cohort.

Pair	N	Accuracy ρ	Cost ρ	ND Jaccard
CORE-Bench Hard–GAIA	14	0.60 [0.16, 0.81]	0.93 [0.72, 0.99]	0.43
CORE-Bench Hard–SWE-bench Mini	14	0.48 [-0.05, 0.82]	0.74 [0.29, 0.93]	0.25
CORE-Bench Hard–TAU Airline	12	0.50 [-0.12, 0.86]	0.78 [0.26, 0.98]	0.17
GAIA–SWE-bench Mini	17	0.71 [0.31, 0.91]	0.89 [0.61, 0.99]	0.33
GAIA–TAU Airline	14	0.58 [0.14, 0.85]	0.86 [0.58, 0.97]	0.29
SWE-bench Mini–TAU Airline	14	0.76 [0.40, 0.91]	0.82 [0.46, 0.97]	0.43

2025]. Ndzomga [2026] study within-benchmark rank preservation under task reduction, while Benchmark² studies cross-benchmark rank consistency for static LLMs [Qian et al., 2026]. Open-SWE-Traces releases unencrypted SWE-agent and OpenHands trajectory corpora [Ahmad et al., 2026]. None directly establishes cross-workload rank transfer for interactive-agent dollar costs under a repeated scaffold.

3 Data and method

We analyze the public `all_leaderboards_costs_HAL.csv` snapshot distributed with Ndzomga [2026]: 242 displayed evaluations across nine benchmarks, 11 scaffolds, and 29 model labels. Accuracy and total dollar cost are complete, but 235 rows report one run. The frozen source, exact hash, and code are in the anonymous supplement.

The primary cohort fixes the displayed HAL Generalist Agent scaffold: 21 labels on CORE-Bench Hard, 17 on GAIA, 9 on SciCode, 7 on ScienceAgentBench, 18 on SWE-bench Verified Mini, and 14 on TAU-bench Airline. USACO has one Generalist result and is excluded. For each pair with at least five shared labels, we compute Spearman ρ and Kendall rank correlation for accuracy and dollar cost. We use 10,000 configuration-resampling draws and 10,000 label-pairing permutations; these quantify finite-label sensitivity, not rollout-level or population-model uncertainty. Holm–Bonferroni correction is applied separately to 15 accuracy and 15 cost tests.

A nondominated point has no competitor that is at least as accurate and no more costly with one strict improvement. We also implement a HAL-inspired origin-anchored convex-hull reconstruction and a 5% tolerance sensitivity where dominance requires 5% lower cost and 5% higher accuracy. The tolerance is a near-tie stress test, not a calibrated measurement-error model.

4 Results

4.1 Raw cost-rank transfer is explained by stable label cost propensity

The robust raw correlations invite a confound check: cost rankings may encode persistent displayed-label price/cost tiers rather than portable execution dynamics. The frozen CSV lacks token components, provider alias, and historical price-card fields, so we cannot claim a list-price adjustment. We fit an observational pooled log-cost model on 86 Generalist rows with benchmark and displayed-label effects. It explains 82.0 percent of log-cost variance, with adjusted R-squared of 72.7 percent. The six high-overlap raw correlations collapse after residualization: none of the residual-cost rank correlations is significant under a 20,000-draw permutation test. Thus raw cost transfer largely reflects stable observed label-level cost propensity. This propensity may combine provider price tier, reasoning configuration, token use, cache treatment, routing, and unrecorded setup; it is not identifiable as list price from the public CSV. Exact adjusted coefficients are reported in the anonymous supplement. The raw cross-workload rank-transfer matrix, including each pair’s Spearman correlation and shared-label count, is provided in the anonymous supplement. The raw dollar-cost structure is largely absorbed by stable displayed-label cost propensity.

4.2 Success adjustment reduces portability

Raw dollar cost does not answer which label buys expected success most cheaply. With the pre-specified one-percent denominator floor, mean CostPerSuccess rank correlation across the six high-overlap pairs is 0.59, versus 0.84 for raw cost; five of six pairs decline. The largest changes occur for GAIA–SWE-mini (0.89 to 0.29), CORE–SWE-mini (0.74 to 0.32), and SWE-mini–TAU (0.82 to 0.45); CORE–TAU is the exception. Alternative zero-exclusion and 5




Figure 1: Nondominated cost–accuracy membership across primary workloads. The final column is each label’s rate over available workloads. No broadly tested label appears on every available frontier.

4.3 Held-out portability is benchmark-specific

In leave-one-benchmark-out prediction, displayed-label cost propensity learned elsewhere predicts raw cost ranks for CORE ($\rho = 0.78$, $N = 14$, $p = 0.0016$), GAIA (0.87 , $N = 16$, $p < 0.0001$), SWE-mini (0.90 , $N = 16$, $p < 0.0001$), and TAU (0.83 , $N = 14$, $p = 0.0004$); it does not predict SciCode (-0.17 , $N = 9$) or ScienceAgentBench (-0.64 , $N = 7$). This is benchmark-specific heterogeneity, not a domain claim. The full plot is in the supplement.

4.4 Frontier membership is non-universal

No label tested on at least five workloads lies on every cost–accuracy frontier. This conclusion survives nondominance, the HAL-inspired hull, and 5% near-tie tolerance. The stronger statement that no label appears on more than 4/6 frontiers does not survive: under tolerance, o4-mini Low appears on 5/6. We therefore retain only the robust non-universal-winner conclusion.

4.5 Scaffold and frontier definitions matter

The source HAL labels recover 30 rows under the hull reconstruction, with two discrepancies; ordinary nondominance includes 19 additional points. A discrete buyer choosing one system and a deployer willing to randomize between systems therefore face different decision frontiers. Within individual benchmarks, scaffolds also differ substantially: across 15 shared labels on SWE-mini, displayed SWE-Agent averages 27.4 percentage points higher accuracy than the Generalist scaffold, with higher median cost. These contrasts motivate fixing the scaffold; they are not a general causal scaffold estimate.

5 External matched trace extension

To assess whether a scaffold can alter measurable execution structure independently of the model and task, we analyze a separate public, unencrypted Open-SWE-Traces boundary sample (four documented MiniMax-M2.5 shards, not the full corpus). We exactly match 977 task–model pairs evaluated under SWE-agent and OpenHands. The trace schema provides success labels, structured turns, tool calls, and message text, but not billed tokens

or dollars; we therefore report trace-burden proxies, not cost. OpenHands has lower matched medians for trajectory turns (123 versus 149), assistant turns/tool calls (61 versus 74), and trajectory characters (176,708 versus 194,263); paired Wilcoxon tests are significant for these proxies (all $p \leq 2.0 \times 10^{-12}$). Matched resolved status shows no detectable difference ($p = 0.197$). This external result demonstrates that scaffold choice can change observable trajectory structure without a detected success change; it does not identify the mechanism behind HAL dollar cost or validate dollar-cost effects.

A single benchmark can supply a useful prior about *raw* observed dollar cost for some held-out workloads, but the analysis does not establish portable workload-specific execution-cost dynamics: raw cost-rank transfer is largely absorbed by stable displayed-label cost propensity, and CostPerSuccess is less portable. Because this propensity may include provider price tier, reasoning configuration, token volume, cache treatment, routing, and unrecorded setup, it cannot be reduced to a list-price claim from the frozen public CSV. Practitioners should validate success-adjusted cost and token behavior on representative target tasks rather than treating one leaderboard frontier as a universal efficiency mechanism.

Construct validity is limited because public `Models` and `Agent Name` strings do not fully encode prompts, tools, budgets, harness versions, or pricing assumptions. The CSV also lacks provider aliases and token components, precluding historical list-price or common-tariff reconstruction; public HAL trace archives are encrypted. CostPerSuccess is an aggregate dollars-per-expected-success proxy, not task-level successful-run cost; the 1

6 Conclusion

Raw dollar-cost ranks transfer robustly across several high-overlap public HAL workload pairs, but pooled label and benchmark cost propensity explains most of this signal. Aggregate CostPerSuccess is less portable, showing that predictable cheapness is not equivalent to predictable dollars per expected success. An external exact task/model trace comparison further shows that scaffolds can materially change trajectory turns and tool calls without a detected success difference. The public evidence does not identify dollar-cost mechanisms, tokens, or provider prices, but it demonstrates why raw cost leaderboards combine stable label propensity with execution structure and should be validated with success-adjusted, trace-aware measures on target tasks.

References

- Wasi Uddin Ahmad, Nikolai Ludwig, Somshubra Majumdar, and Boris Ginsburg. Open-SWE-traces: Advancing dual-mode multilingual distillation for software engineering agents. *arXiv preprint arXiv:2606.16038*, 2026. URL <https://arxiv.org/abs/2606.16038>.
- Zhiyu Fan, Kirill Vasilevski, Dayi Lin, et al. SWE-effi: Re-evaluating software AI agent system effectiveness under resource constraints. *arXiv preprint arXiv:2509.09853*, 2025. URL <https://arxiv.org/abs/2509.09853>.
- Sayash Kapoor et al. AI agents that matter. *arXiv preprint arXiv:2407.01502*, 2024. URL <https://arxiv.org/abs/2407.01502>.
- Sayash Kapoor et al. Holistic agent leaderboard: The missing infrastructure for AI agent evaluation. *arXiv preprint arXiv:2510.11977*, 2026. URL <https://arxiv.org/abs/2510.11977>.
- Franck Ndzomga. Efficient benchmarking of AI agents. *arXiv preprint arXiv:2603.23749*, 2026. URL <https://arxiv.org/abs/2603.23749>.
- Qi Qian, Chengsong Huang, et al. Benchmark²: Systematic evaluation of LLM benchmarks. *arXiv preprint arXiv:2601.03986*, 2026. URL <https://arxiv.org/abs/2601.03986>.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction state that the analysis is descriptive, restricted to repeated public HAL display labels under a repeated displayed scaffold, and does not identify causal model or scaffold effects.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

151 Answer: [Yes]
 152 Justification: Section 6 discusses sparse coverage, display-label ambiguity, single-run evaluations,
 153 price dependence, missing token totals, limited domain comparisons, and frontier sensitivity.

154 **3. Theory assumptions and proofs**
 155 Question: For each theoretical result, does the paper provide the full set of assumptions and a complete
 156 (and correct) proof?
 157 Answer: [N/A]
 158 Justification: The paper presents an empirical secondary analysis and no theoretical results.

159 **4. Experimental result reproducibility**
 160 Question: Does the paper fully disclose all the information needed to reproduce the main experimental
 161 results of the paper to the extent that it affects the main claims and/or conclusions of the paper
 162 (regardless of whether the code and data are provided or not)?
 163 Answer: [Yes]
 164 Justification: Sections 3–4 and the appendix specify the public source, frozen hash, inclusion criteria,
 165 matching unit, frontier definitions, overlap threshold, resampling procedure, and seed.

166 **5. Open access to data and code**
 167 Question: Does the paper provide open access to the data and code, with sufficient instructions to
 168 faithfully reproduce the main experimental results, as described in supplemental material?
 169 Answer: [Yes]
 170 Justification: An anonymized supplement will contain the analysis code, provenance manifest, and
 171 reproduction instructions; the underlying public data source is cited.

172 **6. Experimental setting/details**
 173 Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters,
 174 how they were chosen, type of optimizer) necessary to understand the results?
 175 Answer: [Yes]
 176 Justification: No models are trained. The data filtering, rank metrics, frontier definitions, resampling
 177 count, and seed are specified in Sections 3–4.

178 **7. Experiment statistical significance**
 179 Question: Does the paper report error bars suitably and correctly defined or other appropriate informa-
 180 tion about the statistical significance of the experiments?
 181 Answer: [Yes]
 182 Justification: We report configuration-resampling sensitivity intervals and explicitly state that they are
 183 not rollout-level or population-model confidence intervals and are not used for hypothesis tests.

184 **8. Experiments compute resources**
 185 Question: For each experiment, does the paper provide sufficient information on the computer
 186 resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?
 187 Answer: [N/A]
 188 Justification: The work is a lightweight secondary analysis of public tabular data and uses no model
 189 training or agent rollouts.

190 **9. Code of ethics**
 191 Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code
 192 of Ethics <https://neurips.cc/public/EthicsGuidelines>?
 193 Answer: [Yes]
 194 Justification: The work uses public aggregate evaluation data, attributes source authors, reports
 195 limitations, and makes no claims about people or protected groups.

196 **10. Broader impacts**
 197 Question: Does the paper discuss both potential positive societal impacts and negative societal impacts
 198 of the work performed?
 199 Answer: [N/A]
 200 Justification: The work is a methodological analysis of public leaderboard measurements. Its practical
 201 implication is limited to more cautious agent selection and evaluation reporting.

202 **11. Safeguards**

203 Question: Does the paper describe safeguards that have been put in place for responsible release of
 204 data or models that have a high risk for misuse (e.g., pre-trained language models, image generators,
 205 or scraped datasets)?
 206 Answer: [N/A]
 207 Justification: The paper releases no trained model, agent capability, or new high-risk dataset.

208 **12. Licenses for existing assets**
 209 Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper,
 210 properly credited and are the license and terms of use explicitly mentioned and properly respected?
 211 Answer: [Yes]
 212 Justification: The data and methodological sources are cited; the supplement records the upstream
 213 repository, commit, source path, and declared MIT license.

214 **13. New assets**
 215 Question: Are new assets introduced in the paper well documented and is the documentation provided
 216 alongside the assets?
 217 Answer: [Yes]
 218 Justification: The analysis code and generated tables are documented by a source manifest, artifact
 219 manifest, tests, and reproduction commands.

220 **14. Crowdsourcing and research with human subjects**
 221 Question: For crowdsourcing experiments and research with human subjects, does the paper include
 222 the full text of instructions given to participants and screenshots, if applicable, as well as details about
 223 compensation (if any)?
 224 Answer: [N/A]
 225 Justification: The study involves no crowdsourcing and no human subjects research.

226 **15. Institutional review board (IRB) approvals or equivalent for research with human subjects**
 227 Question: Does the paper describe potential risks incurred by study participants, whether such
 228 risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an
 229 equivalent approval/review based on the requirements of your country or institution) were obtained?
 230 Answer: [N/A]
 231 Justification: The study involves no human subjects research.

232 **16. Declaration of LLM usage**
 233 Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard
 234 component of the core methods in this research?
 235 Answer: [N/A]
 236 Justification: The core method is deterministic analysis of public tabular data; no LLM is used as a
 237 component of the empirical method.