
Does Cost-Efficiency Travel?

What Agent Leaderboards Can Predict About Cost Across Workloads

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agent leaderboards increasingly report task success and dollar cost, inviting users to choose
2 a “cost-efficient” system from a single benchmark. We ask whether that recommendation
3 travels to a new workload. Using public Holistic Agent Leaderboard (HAL) results, we hold
4 the displayed HAL Generalist Agent scaffold fixed and compare repeated displayed model
5 labels across six workloads. Raw dollar-cost ranks are strongly associated in several high-
6 overlap pairs, but this signal is largely absorbed by stable displayed-label and benchmark
7 cost propensity. Held-out prediction is strong for four workloads and uninformative for
8 two lower-overlap workloads. Aggregate dollars per expected success are less stable: only
9 three CostPerSuccess associations survive permutation testing and Holm correction, and
10 the clearest raw-cost versus CostPerSuccess gap occurs for GAIA-SWE-mini. A separate
11 matched public trace study gives the system-level context that leaderboard aggregates lack:
12 across 977 exact task–model pairs, OpenHands uses about 15% fewer median turns and
13 tool calls than SWE-agent, without a detected difference in resolved status. Taken together,
14 the results show that raw cost is a useful but incomplete portability signal. It combines
15 persistent label-level cost propensity with scaffold execution structure, and it should not be
16 treated as a portable ranking of deployment efficiency.

17 1 The practical problem

18 Agent evaluation is expensive enough that cost has become a first-class leaderboard number. HAL reports
19 21,730 rollouts across nine benchmarks and roughly \$40,000 in evaluation cost [Kapoor et al., 2026]. That
20 makes a cost–accuracy frontier appealing: if two systems perform similarly, why not choose the cheaper one?

21 The hidden step in that reasoning is portability. A user sees that a label is inexpensive on one benchmark
22 and assumes it will remain a good buy in a different environment. That assumption is plausible, but it is not
23 innocuous. A reported dollar cost can reflect the model or provider, reasoning mode, prompt and tool setup,
24 cache behavior, scaffold, and the workload itself.

25 The GAIA-SWE-mini comparison makes the issue concrete. Across 17 shared displayed labels, raw dollar-cost
26 ranks are closely aligned (Spearman $\rho = 0.89$). But the ranking by aggregate dollars per expected success is
27 much less aligned ($\rho = 0.29$); the paired raw-minus-CPS contrast is 0.60 with a 95% configuration-resampling
28 sensitivity interval of [0.14, 1.13]. In practical terms, a GAIA-like leaderboard can offer a useful prior about
29 raw spend for a SWE-like workload without providing a reliable ordering of expected dollars per successful
30 task. This is an aggregate ranking example, not a claim about a named model or a production deployment.

31 We therefore ask three questions. First, when does an observed cost ranking predict a held-out workload? Sec-
32 ond, does that ranking still travel after separating stable displayed-label cost propensity from benchmark cost
33 level? Third, what can public traces say about the scaffold-level execution structure that aggregate leaderboards
34 hide?

2 Why this is different from prior cost and benchmark work

This paper is an empirical evaluation study, not a new rank estimator. Kapoor et al. [2024] established the importance of reporting accuracy and cost jointly and motivates the frontier perspective used here. Ndzomga [2026] asks whether a smaller task subset preserves agent rankings within a benchmark; we instead ask whether an observed cost signal predicts a different workload. Benchmark² studies cross-benchmark rank consistency for static LLM benchmarks; our setting adds interactive scaffolds, dollar cost, success adjustment, and a held-out workload test. SWE-Effi examines resource effectiveness within software engineering; our result concerns the portability of public leaderboard signals across workloads. The matched Open-SWE extension does not claim to reproduce HAL cost measurements. It supplies separate, exact task–model evidence that scaffold choice can change observable trajectory burden.

Our primary data are the public `all_leaderboards_costs_HAL.csv` snapshot distributed with Ndzomga [2026]: 242 displayed evaluations across nine benchmarks, 11 scaffolds, and 29 displayed model labels. The coverage determines the design. Most scaffolds appear on one workload only, so the data do not support a general model–scaffold–workload decomposition. We instead fix the displayed HAL Generalist Agent scaffold, match exact public model-label strings, and retain six workloads with enough repeated labels: CORE-Bench Hard, GAIA, SciCode, ScienceAgentBench, SWE-bench Verified Mini, and TAU-bench Airline. This cohort contains 86 rows. USACO has one Generalist observation and is excluded.

The primary CSV supplies accuracy and aggregate dollar cost, but not token components, provider aliases, API calls, tool calls, retries, cache usage, or per-task trajectories. Accordingly, we do not call any adjustment “price controlled,” and we do not infer token or failure costs from dollars. Most evaluations are single-run, so rollout-level uncertainty is unavailable.

For each pair with at least five shared labels, we rank the same shared displayed labels by each metric. A configuration-resampling sensitivity interval is obtained by drawing N shared labels with replacement, keeping each draw’s paired values together, and recomputing the statistic; 10,000 draws yield its 2.5th and 97.5th percentiles. It describes sensitivity to the finite observed label set, not a rollout-level confidence interval. Raw-cost Holm correction applies to 15 tests. CPS has its own 20,000-draw label-pairing permutation tests and Holm correction over 15 tests.

$$\text{CPS}_{i,b}(\epsilon) = \frac{C_{i,b}}{\max(A_{i,b}, \epsilon)},$$

where A is accuracy as a fraction and the primary floor is $\epsilon = 0.01$. The 5% and zero-exclusion versions are reported in the supplement. CPS is an aggregate dollars-per-expected-success proxy, not the mean cost of successful runs.

To test out-of-sample portability, we use leave-one-benchmark-out (LOBO) prediction. For each held-out workload, we estimate each label’s centered log-cost propensity from the other workloads and predict the held-out cost rank. We also fit the descriptive pooled model

$$\log C_{i,b} = \mu + \alpha_i + \beta_b + \epsilon_{i,b}.$$

This 86-row model has 30 parameters and 56 residual degrees of freedom. Its fit is descriptive; LOBO is the stronger test of portability.

Finally, we report three decision frontiers: discrete nondominance, a HAL-inspired origin-anchored convex hull, and a 5% near-tie tolerance. These answer related but different decision questions.

3 Evidence

3.1 Raw cost predicts some held-out workloads

The most direct evidence comes from LOBO. A label’s cost propensity learned from the other workloads predicts held-out raw cost ranks for CORE ($\rho = 0.78$, $N = 14$, $p = 0.0016$), GAIA (0.87 , $N = 16$, $p < 0.0001$), SWE-mini (0.90 , $N = 16$, $p < 0.0001$), and TAU (0.83 , $N = 14$, $p = 0.0004$). These are not one-label artifacts: leave-one-label-out ranges remain positive for every omission, from $[0.73, 0.91]$ for CORE to $[0.88, 0.92]$ for SWE-mini.

The picture is not uniform. SciCode is unstable under leave-one-label-out analysis ($[-0.50, 0.14]$). ScienceAgentBench remains directionally negative ($[-0.77, -0.43]$), but it has seven labels and its original permutation result is not significant. SciCode has very low mean accuracy (2.1%) and one-third zero-accuracy labels; ScienceAgentBench has the narrowest log-cost dispersion. These are useful diagnostic facts, not a basis for a domain-level causal story.

84 3.2 What the raw-cost signal appears to contain

85 Among high-overlap pairs, raw dollar-cost correlations range from 0.74 to 0.93. A pooled benchmark-plus-
86 label log-cost model has $R^2 = 82.0\%$ and adjusted $R^2 = 72.7\%$, but this in-sample fit should not be overread
87 given 30 parameters and 86 rows. Its purpose is diagnostic. Once observed benchmark and displayed-label cost
88 propensity are removed, none of the high-overlap residual rank associations survives a 20,000-draw permutation
89 test.

90 The reasonable interpretation is not that cost is meaningless. Rather, much of the portable raw-cost signal
91 is a stable observed label-level component. The public data cannot identify whether that component comes
92 from provider pricing, reasoning settings, routing, token behavior, cache treatment, or other unrecorded setup
93 choices.

94 3.3 Cheapness is not the same ranking as dollars per expected success

95 CPS changes the question from “which label is cheap?” to “which label buys expected success cheaply?” Under
96 the 1% floor, CPS association survives 20,000-draw permutation testing and Holm correction for CORE–GAIA
97 ($\rho = 0.87$, 95% sensitivity interval [0.63, 0.96]), CORE–TAU (0.83, [0.31, 1.00]), and GAIA–TAU (0.78, [0.42,
98 0.94]). Most other CPS associations remain uncertain.

99 The high-overlap mean CPS correlation is 0.59, compared with 0.84 for raw cost, but that average is descriptive
100 rather than a population-level contrast. The clearest paired difference is GAIA–SWE-mini: raw cost exceeds
101 CPS correlation by 0.60, with a 95% resampling interval [0.14, 1.13]. In short, a raw-spend ranking can travel
102 while a success-adjusted ranking does not.

103 3.4 There is no universal cost–accuracy winner

104 No broadly tested label lies on every frontier under nondominance, the HAL-inspired hull, or the 5% tolerance
105 rule. This is a practical guardrail, not the paper’s central novelty. The stronger statement that no label appears
106 on more than 4/6 frontiers fails under tolerance: o4-mini Low appears on 5/6.

107 4 A separate matched trace study

108 The leaderboard analysis tells us what travels, but not what happens inside an agent run. HAL public trace
109 archives are encrypted. We therefore use a separate public Open-SWE-Traces boundary sample: four docu-
110 mented MiniMax-M2.5 shards under SWE-agent and OpenHands. It is not a random sample of the full corpus.
111 We exactly match task ID and model, exclude duplicate task/model/scaffold records, and obtain 977 pairs.

112 The traces expose resolved status, structured messages, tool calls, and text. They do not expose billing, tokens,
113 provider usage, retries, or latency. We therefore call turns, tool calls, and message characters trace-burden
114 proxies, not cost.

115 At fixed task and model, OpenHands/SWE-agent has a median ratio of 0.85 for trajectory turns (10,000-
116 bootstrap interval [0.82, 0.87]) and 0.85 for tool calls ([0.82, 0.86]). OpenHands has lower values on 69.7% of
117 pairs for each measure. Total trajectory characters have a smaller ratio effect, 0.95 ([0.92, 0.98]), with Open-
118 Hands lower on 55.1% of pairs. Resolved status has no detected paired difference ($p = 0.197$). Thus, scaffolds
119 can change visible execution burden even when task and model are held fixed. This does not show lower token
120 use or dollar cost, and it does not identify the mechanism behind HAL rankings.

121 5 Interpretation and implications

122 The evidence supports a narrower conclusion than a single leaderboard recommendation suggests. Raw dollar
123 cost can be predictive across some workloads. In the HAL cohort, that prediction is robust for four held-out
124 workloads and is not driven by a single displayed label. But the signal is partly stable label-level propensity, not
125 a demonstrated measure of portable execution efficiency. Success adjustment changes the ranking and is only
126 partly portable. The matched trace study supplies a compatible systems-level explanation: scaffold choices
127 alter observable work even when task and model are fixed.

128 For practitioners, this means that a leaderboard’s dollar figure is useful as a first screen, not as a final procure-
129 ment decision. A target evaluation should include a success-adjusted metric and, where possible, trajectory
130 telemetry. For benchmark builders, the missing fields are concrete: model/provider aliases, input/output/cache
131 usage, calls, tool activity, retries, outcome labels, and per-task costs would make it possible to distinguish price
132 from execution behavior.

6 Limits

This is a secondary-data analysis. Public display labels do not fully identify prompts, tools, budgets, harness versions, or routing. The HAL cohort is sparse and unbalanced, with mostly single-run rows. The public HAL snapshot cannot identify failure-cost ratios, retry behavior, token trajectories, tail cost, or a causal model–scaffold–workload decomposition. CPS is not task-level successful-run cost. The Open-SWE extension is external to HAL and uses a reproducible four-shard boundary sample, not the full corpus. Its proxies are not tokens or dollars. These limits define the claims rather than invalidate them.

7 Conclusion

A cost-aware leaderboard is better than an accuracy-only leaderboard, but it is not a portable deployment-efficiency ranking by default. In these public agent data, raw cost can generalize to some held-out workloads, yet it partly reflects stable displayed-label propensity and does not always preserve the ranking by dollars per expected success. Matched traces show that scaffold design changes observable execution burden at fixed task and model. The next generation of agent leaderboards should make these distinctions measurable rather than leaving them hidden inside one dollar number.

References

- Sayash Kapoor et al. AI agents that matter. *arXiv preprint arXiv:2407.01502*, 2024. URL <https://arxiv.org/abs/2407.01502>.
- Sayash Kapoor et al. Holistic agent leaderboard: The missing infrastructure for AI agent evaluation. *arXiv preprint arXiv:2510.11977*, 2026. URL <https://arxiv.org/abs/2510.11977>.
- Franck Ndzomga. Efficient benchmarking of AI agents. *arXiv preprint arXiv:2603.23749*, 2026. URL <https://arxiv.org/abs/2603.23749>.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state that the analysis is descriptive, restricted to repeated public HAL display labels under a repeated displayed scaffold, and does not identify causal model or scaffold effects.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 discusses sparse coverage, display-label ambiguity, single-run evaluations, price dependence, missing token totals, limited domain comparisons, and frontier sensitivity.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper presents an empirical secondary analysis and no theoretical results.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Sections 3–4 and the appendix specify the public source, frozen hash, inclusion criteria, matching unit, frontier definitions, overlap threshold, resampling procedure, and seed.

5. Open access to data and code

180 Question: Does the paper provide open access to the data and code, with sufficient instructions to
181 faithfully reproduce the main experimental results, as described in supplemental material?

182 Answer: [Yes]

183 Justification: An anonymized supplement will contain the analysis code, provenance manifest, and
184 reproduction instructions; the underlying public data source is cited.

185 **6. Experimental setting/details**

186 Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters,
187 how they were chosen, type of optimizer) necessary to understand the results?

188 Answer: [Yes]

189 Justification: No models are trained. The data filtering, rank metrics, frontier definitions, resampling
190 count, and seed are specified in Sections 3–4.

191 **7. Experiment statistical significance**

192 Question: Does the paper report error bars suitably and correctly defined or other appropriate infor-
193 mation about the statistical significance of the experiments?

194 Answer: [Yes]

195 Justification: We report configuration-resampling sensitivity intervals and explicitly state that they
196 are not rollout-level or population-model confidence intervals and are not used for hypothesis tests.

197 **8. Experiments compute resources**

198 Question: For each experiment, does the paper provide sufficient information on the computer re-
199 sources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

200 Answer: [N/A]

201 Justification: The work is a lightweight secondary analysis of public tabular data and uses no model
202 training or agent rollouts.

203 **9. Code of ethics**

204 Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS
205 Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

206 Answer: [Yes]

207 Justification: The work uses public aggregate evaluation data, attributes source authors, reports limi-
208 tations, and makes no claims about people or protected groups.

209 **10. Broader impacts**

210 Question: Does the paper discuss both potential positive societal impacts and negative societal im-
211 pacts of the work performed?

212 Answer: [Yes]

213 Justification: We use ****Yes**** here because the work has practical consequences. Better reporting of
214 success-adjusted cost and execution telemetry can reduce wasted evaluation and deployment spend.
215 Conversely, treating a single leaderboard’s dollar ranking as a deployment recommendation can mis-
216 allocate resources or favor systems whose price/configuration signal does not transfer. The paper’s
217 safeguards are explicit claim boundaries, public-data provenance, and reproducible sensitivity analy-
218 ses.

219 **11. Safeguards**

220 Question: Does the paper describe safeguards that have been put in place for responsible release of
221 data or models that have a high risk for misuse (e.g., pre-trained language models, image generators,
222 or scraped datasets)?

223 Answer: [N/A]

224 Justification: The paper releases no trained model, agent capability, or new high-risk dataset.

225 **12. Licenses for existing assets**

226 Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper,
227 properly credited and are the license and terms of use explicitly mentioned and properly respected?

228 Answer: [Yes]

229 Justification: The data and methodological sources are cited; the supplement records the upstream
230 repository, commit, source path, and declared MIT license.

231 **13. New assets**

232 Question: Are new assets introduced in the paper well documented and is the documentation provided
 233 alongside the assets?

234 Answer: [Yes]

235 Justification: The analysis code and generated tables are documented by a source manifest, artifact
 236 manifest, tests, and reproduction commands.

237 **14. Crowdsourcing and research with human subjects**

238 Question: For crowdsourcing experiments and research with human subjects, does the paper include
 239 the full text of instructions given to participants and screenshots, if applicable, as well as details about
 240 compensation (if any)?

241 Answer: [N/A]

242 Justification: The study involves no crowdsourcing and no human subjects research.

243 **15. Institutional review board (IRB) approvals or equivalent for research with human subjects**

244 Question: Does the paper describe potential risks incurred by study participants, whether such risks
 245 were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equiv-
 246 alent approval/review based on the requirements of your country or institution) were obtained?

247 Answer: [N/A]

248 Justification: The study involves no human subjects research.

249 **16. Declaration of LLM usage**

250 Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard
 251 component of the core methods in this research?

252 Answer: [N/A]

253 Justification: The core method is deterministic analysis of public tabular data; no LLM is used as a
 254 component of the empirical method.